

## Explore the data

Our modeling goal here is to predict which countries are Asian countries and which countries are not, based on their patterns of food consumption in the eleven categories from the [#TidyTuesday dataset](#). The original data is in a long, tidy format, and includes information on the carbon emission associated with each category of food consumption.

```
library(tidyverse)

food_consumption <- readr::read_csv("https://raw.githubusercontent.com/
rfordatascience/tidyuesday/master/data/2020/2020-02-18/food_consumption.csv")

food_consumption

## # A tibble: 1,430 x 4
##   country    food_category consumption co2_emmission
##
## 1 Argentina Pork                10.5          37.2
## 2 Argentina Poultry             38.7          41.5
## 3 Argentina Beef                55.5         1712
## 4 Argentina Lamb & Goat         1.56          54.6
## 5 Argentina Fish                 4.36           6.96
## 6 Argentina Eggs                11.4          10.5
## 7 Argentina Milk - inc. cheese  195.          278.
## 8 Argentina Wheat and Wheat Products 103.          19.7
## 9 Argentina Rice                 8.77          11.2
## 10 Argentina Soybeans            0              0
## # ... with 1,420 more rows
```

Let's build a dataset for modeling that is wide instead of long using `pivot_wider()` from `tidyr`. We can use the [countrycode](#) package to find which continent each country is in, and create a new variable for prediction `asia` that tells us whether a country is in Asia or not.

```
library(countrycode)
library(janitor)

food <- food_consumption %>%
  select(-co2_emmission) %>%
  pivot_wider(
    names_from = food_category,
    values_from = consumption
  ) %>%
  clean_names() %>%
  mutate(continent = countrycode(
    country,
    origin = "country.name",
    destination = "continent"
  )) %>%
  mutate(asia = case_when(
    continent == "Asia" ~ "Asia",
    TRUE ~ "Other"
  )) %>%
  select(-country, -continent) %>%
  mutate_if(is.character, factor)

food
```

```
## # A tibble: 130 x 12
##   pork poultry beef lamb_goat fish eggs milk_inc_cheese wheat_and_wheat...
##
## 1 10.5      38.7 55.5      1.56 4.36 11.4              195.          103.
## 2 24.1      46.1 33.9      9.87 17.7 8.51              234.          70.5
## 3 10.9      13.2 22.5     15.3 3.85 12.5              304.          139.
## 4 21.7      26.9 13.4     21.1 74.4 8.24              226.          72.9
## 5 22.3      35.0 22.5     18.9 20.4 9.91              137.          76.9
## 6 27.6      50.0 36.2      0.43 12.4 14.6              255.          80.4
## 7 16.8      27.4 29.1      8.23 6.53 13.1              211.          109.
## 8 43.6      21.4 29.9      1.67 23.1 14.6              255.          103.
## 9 12.6      45    39.2      0.62 10.0 8.98              149.           53
## 10 10.4     18.4 23.4      9.56 5.21 8.29              288.          92.3
## # ... with 120 more rows, and 4 more variables: rice , soybeans ,
## #   nuts_inc_peanut_butter , asia
```

This is not a big dataset, but it will be good for demonstrating how to tune hyperparameters. Before we get started on that, how are the categories of food consumption related? Since these are all numeric variables, we can use `ggscatmat()` for a quick visualization.

```
library(GGally)
ggscatmat(food, columns = 1:11, color = "asia", alpha = 0.7)
```



Notice how important rice is! Also see how the relationships between different food categories is different for Asian and non-Asian countries; a tree-based model like a random forest is good at learning interactions like this.

## Tune hyperparameters

Now it's time to tune the hyperparameters for a random forest model. First, let's create a set of bootstrap resamples to use for tuning, and then let's create a model specification for a random forest where we will tune `mtry` (the number of predictors to sample at each split) and `min_n` (the number of observations needed to keep splitting nodes). There are **hyperparameters** that can't be learned from data when training the model.

```
library(tidymodels)

set.seed(1234)
food_boot <- bootstraps(food, times = 30)
food_boot

## # Bootstrap sampling
## # A tibble: 30 x 2
##   splits          id
##
## 1 Bootstrap01
## 2 Bootstrap02
## 3 Bootstrap03
## 4 Bootstrap04
## 5 Bootstrap05
## 6 Bootstrap06
## 7 Bootstrap07
## 8 Bootstrap08
## 9 Bootstrap09
## 10 Bootstrap10
## # ... with 20 more rows
```

```
rf_spec <- rand_forest(
  mode = "classification",
  mtry = tune(),
  trees = 1000,
  min_n = tune()
) %>%
  set_engine("ranger")

rf_spec

## Random Forest Model Specification (classification)
##
## Main Arguments:
##   mtry = tune()
##   trees = 1000
##   min_n = tune()
##
## Computational engine: ranger
```

We can't learn the right values when training a single model, but we can train a whole bunch of models and see which ones turn out best. We can use parallel processing to make this go faster, since the different parts of the grid are independent.

```
doParallel::registerDoParallel()

rf_grid <- tune_grid(
  asia ~ .,
  model = rf_spec,
  resamples = food_boot
)

rf_grid

## # Bootstrap sampling
## # A tibble: 30 x 4
##   splits          id      .metrics      .notes
## *
## 1 Bootstrap01
## 2 Bootstrap02
## 3 Bootstrap03
## 4 Bootstrap04
## 5 Bootstrap05
## 6 Bootstrap06
## 7 Bootstrap07
## 8 Bootstrap08
## 9 Bootstrap09
## 10 Bootstrap10
## # ... with 20 more rows
```

Once we have our tuning results, we can check them out.

```
rf_grid %>%
  collect_metrics()

## # A tibble: 20 x 7
##   mtry min_n .metric .estimator mean      n std_err
##
## 1     2     4 accuracy binary    0.836  30 0.00798
## 2     2     4 roc_auc  binary    0.843  30 0.00861
```

```
## 3      2      12 accuracy binary    0.830    30 0.00760
## 4      2      12 roc_auc  binary    0.833    30 0.00930
## 5      4      33 accuracy binary    0.815    30 0.00873
## 6      4      33 roc_auc  binary    0.818    30 0.0101
## 7      4      37 accuracy binary    0.814    30 0.00875
## 8      4      37 roc_auc  binary    0.820    30 0.0103
## 9      5      31 accuracy binary    0.817    30 0.00864
## 10     5      31 roc_auc  binary    0.822    30 0.0103
## 11     6       9 accuracy binary    0.824    30 0.00895
## 12     6       9 roc_auc  binary    0.831    30 0.00947
## 13     7      21 accuracy binary    0.815    30 0.00947
## 14     7      21 roc_auc  binary    0.824    30 0.0101
## 15     8      18 accuracy binary    0.817    30 0.00929
## 16     8      18 roc_auc  binary    0.824    30 0.0103
## 17     9      26 accuracy binary    0.816    30 0.0100
## 18     9      26 roc_auc  binary    0.822    30 0.0104
## 19    11      15 accuracy binary    0.813    30 0.0110
## 20    11      15 roc_auc  binary    0.825    30 0.0102
```

And we can see which models performed the best, in terms of some given metric.

```
rf_grid %>%
  show_best("roc_auc")

## # A tibble: 5 x 7
##   mtry min_n .metric .estimator  mean     n std_err
##
## 1     2      4 roc_auc binary    0.843    30 0.00861
## 2     2     12 roc_auc binary    0.833    30 0.00930
## 3     6      9 roc_auc binary    0.831    30 0.00947
## 4    11     15 roc_auc binary    0.825    30 0.0102
## 5     8     18 roc_auc binary    0.824    30 0.0103
```